

# Momentum Contrast for Unsupervised Visual Representation Learning (CVPR, 2020)

- Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, Ross Girshick

Presentation Author  
MATHIEU RAVAUT


Nanyang Technological University

# Table of content

- 1 Introduction
- 2 Approach
- 3 Results
- 4 Conclusion

# Introduction

- **Contrastive learning** is a paradigm in machine learning which consists in matching together data points which are similar, and pushing apart the ones which are not.
- A **query** is being compared (with a **similarity**) to a set of **k** different **keys**, among which there is a **positive one** and many **negative** ones.



# Introduction

- Contrastive learning does not necessarily need labels.
- Typically : use the rest of the batch as (easy) negatives.
- Here, the authors propose **MoCo**, a new contrastive learning objective to pre-train vision models in an unsupervised fashion.
- It breaks the **supervised pre-training on ImageNet** paradigm in vision.
- MoCo produces image representations which transfer very well on many downstream tasks : ImageNet linear classification protocol, PASCAL VOC, COCO, etc.

# Introduction

Multiple questions and design choices arise :

- 1 How to build the (query, positive key) pair ?  
This is referred to as the **pretext task**.
- 2 How many  $k$  negative samples to use ? **MAJOR CHOICE**
- 3 How to store all these negative samples ? **MAJOR ISSUE**
- 4 Which similarity function to use ?
- 5 Which contrastive loss function to use ?

# Approach

[1] MoCo's pretext task is *instance discrimination* :

- Apply data augmentation on the images : a  $224 \times 224$ -pixel crop is taken from a randomly resized image, and then undergoes random color jittering, random horizontal flip, and random grayscale conversion. This gives us **multiple views of each image**.
- To build the queries and keys : final layer of a **ResNet-50 encoder**, add a few fully-connected layers with **non-linearities (MoCo v2)**, and apply **L2-norm**.
- **Same** ResNet-50 architecture for both queries and keys encoders.
- A (query, key) pair is positive if they are encoded views of *the same image*.

# Approach

[2] & [3] MoCo's main novelty resides in the way it treats keys :

- Store all keys in a **dictionary**.
- This dictionary is a **queue** of data points from the recent batches : enqueue the current batch, dequeue the last one.
- This *decouples the number of negatives  $k$  from the batch size*.

Concretely :

If the batch size is  $n$ , and we use  $k = p * n$ , the keys dictionary is the current batch + the previous  $(p - 1)$  batches :

$$keys = \{x_0, \dots, x_{n-1}, x_n, \dots, x_{2*n-1}, \dots, x_{(p-1)*n}, \dots, x_{p*n-1}\}$$

# Approach

There's a new issue : since the keys  $\{x_n, \dots x_{2*n-1}, \dots, x_{(p-1)*n}, \dots x_{p*n-1}\}$  were encoded at the previous iterations, the **encoder weights were different**, thus the keys representations are **inconsistent** with the representation of the queries and the current batch keys.



# Approach

To alleviate this issue, the authors propose **momentum update** :

- The key encoder  $\theta_k$  has the same architecture as the query encoder  $\theta_q$ , but different weights.
- The key encoder is not updated by gradient descent.
- Instead, it is updated by **momentum updates** :  
$$\theta_k \leftarrow m * \theta_k + (1 - m) * \theta_q$$
- $m$  must be really high (0.999) : slow updates.
- We thus have another neural network than the base model, in parallel with it and with the same architecture, but not updated via gradient descent : the **momentum model**.

# Approach

The base model is updated with the contrastive loss :

$$\mathcal{L}_q = -\log \frac{\exp(q \cdot k_+ / \tau)}{\sum_{i=0}^K \exp(q \cdot k_i / \tau)}$$


# Approach

One last issue : ResNet-50 uses Batch-Norm, but the model appears to learn the pretext task too easily with it.

Thus, MoCo uses **shuffling BatchNorm** : train with multiple GPUs and perform BN on the samples independently for each GPU. They shuffle the sample order in the current mini-batch for the key encoder before distributing it among GPUs, but NOT for the query encoder.


# Approach

Model recap :



# Results

Linear protocol on ImageNet : train a linear classifier on the image representations obtained with the pre-training.

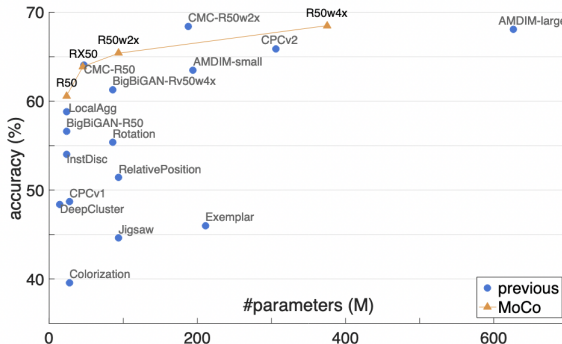


MoCo enables to use a very large k (65,536!), and accuracy keeps increasing : +2.9 point when k goes from 1024 to 16384.



# Results

Comparison with other pre-training methods on the linear protocol on ImageNet :



# Results

Comparison with other pre-training methods on the linear protocol on ImageNet :

method	architecture	#params (M)	accuracy (%)
Exemplar [17]	R50w3×	211	46.0 [38]
RelativePosition [13]	R50w2×	94	51.4 [38]
Jigsaw [45]	R50w2×	94	44.6 [38]
Rotation [19]	Rv50w4×	86	55.4 [38]
Colorization [64]	R101*	28	39.6 [14]
DeepCluster [3]	VGG [53]	15	48.4 [4]
BigBiGAN [16]	R50	24	56.6
	Rv50w4×	86	61.3
<i>methods based on contrastive learning follow:</i>			
InstDisc [61]	R50	24	54.0
LocalAgg [66]	R50	24	58.8
CPC v1 [46]	R101*	28	48.7
CPC v2 [35]	R170* <sub>wider</sub>	303	65.9
CMC [56]	R50 <sub>L+ab</sub>	47	64.1 <sup>†</sup>
	R50w2× <sub>L+ab</sub>	188	68.4 <sup>†</sup>
AMDIM [2]	AMDIM <sub>small</sub>	194	63.5 <sup>†</sup>
	AMDIM <sub>large</sub>	626	68.1 <sup>†</sup>
<b>MoCo</b>	R50	24	60.6
	RX50	46	63.9
	R50w2×	94	65.4
	R50w4×	375	<b>68.6</b>



# Results

On Pascal VOC, MoCo representations are better than supervised pre-training ones :

pre-train	AP <sub>50</sub>	AP	AP <sub>75</sub>
random init.	64.4	37.9	38.6
super. IN-1M	81.4	54.0	59.1
<b>MoCo</b> IN-1M	81.1 (−0.3)	54.6 (+0.6)	59.9 (+0.8)
<b>MoCo</b> IG-1B	81.6 (+0.2)	55.5 (+1.5)	61.2 (+2.1)

(a) Faster R-CNN, R50-dilated-C5

pre-train	AP <sub>50</sub>	AP	AP <sub>75</sub>
random init.	60.2	33.8	33.1
super. IN-1M	81.3	53.5	58.8
<b>MoCo</b> IN-1M	81.5 (+0.2)	55.9 (+2.4)	62.6 (+3.8)
<b>MoCo</b> IG-1B	82.2 (+0.9)	57.2 (+3.7)	63.7 (+4.9)

(b) Faster R-CNN, R50-C4

# Results

MoCo representations are great on multiple vision tasks (object detection on COCO, pose estimation, segmentation on LVIS) :

COCO keypoint detection			
pre-train	AP <sup>kp</sup>	AP <sup>kp</sup> <sub>50</sub>	AP <sup>kp</sup> <sub>75</sub>
random init.	65.9	86.5	71.7
super. IN-1M	65.8	86.9	71.9
<b>MoCo</b> IN-1M	66.8 (+1.0)	87.4 (+0.5)	72.5 (+0.6)
<b>MoCo</b> IG-1B	66.9 (+1.1)	87.8 (+0.9)	73.0 (+1.1)
COCO dense pose estimation			
pre-train	AP <sup>dp</sup>	AP <sup>dp</sup> <sub>50</sub>	AP <sup>dp</sup> <sub>75</sub>
random init.	39.4	78.5	35.1
super. IN-1M	48.3	85.6	50.6
<b>MoCo</b> IN-1M	50.1 (+1.8)	86.8 (+1.2)	53.9 (+3.3)
<b>MoCo</b> IG-1B	50.6 (+2.3)	87.0 (+1.4)	54.3 (+3.7)
LVIS v0.5 instance segmentation			
pre-train	AP <sup>mk</sup>	AP <sup>mk</sup> <sub>50</sub>	AP <sup>mk</sup> <sub>75</sub>
random init.	22.5	34.8	23.8
super. IN-1M <sup>†</sup>	24.4	37.8	25.8
<b>MoCo</b> IN-1M	24.1 (−0.3)	37.4 (−0.4)	25.5 (−0.3)
<b>MoCo</b> IG-1B	24.9 (+0.5)	38.2 (+0.4)	26.4 (+0.6)

# Conclusion

## Takeaways :

- New method for unsupervised pretraining leveraging contrastive learning.
- Enables to scale the number of negatives very high thanks to the queue and the decoupling with the batch size, which drastically improves performance.
- Great performance on multiple vision tasks (classification, detection, etc). Even better than some supervised pre-training methods.
- Consistent gain when going from ImageNet (1M images) to InstaGram (1B images) as pre-training dataset. But still quite a small gain given the data size.

# Conclusion

How can this be applied to NLP :

- I am currently pre-training a new unsupervised summarization method based on MoCo. There are gains from scaling  $k$  from 1024 to 16384.
- In generation tasks in NLP, we have an **encoder-decoder** model, not just a plain **encoder**. Thus, should we use representations from the encoder, decoder, or both ?
- Exploring how to leverage the momentum encoder : use it for data augmentation ?