

Efficiently Modeling Long Sequences with Structured State Spaces (S4)

- Albert Gu, Karan Goel, Christopher Re
(Stanford)

MATHIEU Ravaut

Nanyang Technological University

Overview

- ICLR 2022 with impressive review scores : 8, 8, 8.
- An alternative way of sequence modeling without Transformer or even self-attention.
- Uses **state-space models**.
- Combined with *multiple linear algebra tricks* to keep computation cost low. There'll be *a lot* of maths!
- Extremely good at long sequence modeling : SOTA in Long Range Arena with much faster inference.
- *The Annotated S4*, a dedicated illustrated notebook by Sasha Rush (March 2022).

Table of content

- 1 Introduction
- 2 Model
- 3 Experiments
- 4 Conclusion

State-space models (SSMs)

- State-space models are a classical mathematical tool from the 1970s and 1980s, used for instance in control engineering.
- Framework :
 - A 1-D input signal $u(t)$
 - An N-D latent representation $x(t)$
 - A 1-D output signal $y(t)$
- The state-space model is defined by :

$$x'(t) = \mathbf{A}.x(t) + \mathbf{B}.u(t) \quad (1)$$

$$y(t) = \mathbf{C}.x(t) + \mathbf{D}.u(t) \quad (2)$$

State-space models (SSMs)

- In practice, they ignore **D** because it's simply a skip-connection and set the matrix to 0, so the SSM equation becomes :

$$x'(t) = \mathbf{A}.x(t) + \mathbf{B}.u(t) \quad (3)$$

$$y(t) = \mathbf{C}.x(t) \quad (4)$$

Hippo Matrix

- The basic SSM is a linear first-order ODE.
- It solves with an exponential function.
- Thus, it performs poorly in practice, due to **vanishing/exploding gradient** issues.
- The same authors propose to use the **HiPPO Matrix** as matrix A , from their previous paper *HiPPO : Recurrent memory with optimal polynomial projections* (Neurips 2020).
- This special class of matrices allows the state $x(t)$ to memorize the input $u(t)$ (hard to get the intuition..).

HiPPO Matrix

The HiPPO matrix :

$$\mathbf{A}_{n,k} = - \begin{cases} (2n+1)^{1/2} \cdot (2k+1)^{1/2} & \text{if } n > k \\ (n+1) & \text{if } n = k \\ 0 & \text{if } n < k \end{cases} \quad (5)$$

$$\mathbf{A} = - \begin{bmatrix} 1 & 0 & \dots & 0 & 0 \\ \sqrt{3} & 2 & \dots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \sqrt{2N-1} & \sqrt{(2N-1) \cdot 3} & \dots & \sqrt{(2N-1)(2N-3)} & N \end{bmatrix} \quad (6)$$

Discrete SSM

- So far everything was continuous.
- But in practice, the input signal is discrete $u = (u_0, u_1, \dots)$.
- The *bilinear method* enables to derive a **discrete SSM** :

$$x_k = \overline{\mathbf{A}}.x_{k-1} + \overline{\mathbf{B}}.u_k \quad (7)$$

$$y_k = \overline{\mathbf{C}}.x_k = \mathbf{C}.x_k \quad (8)$$

with the following new matrices (Δ is the step size) :

$$\overline{\mathbf{A}} = (\mathbf{I} - \frac{\Delta}{2}.\mathbf{A})^{-1} . (\mathbf{I} + \frac{\Delta}{2}.\mathbf{A}) \quad (9)$$

$$\overline{\mathbf{B}} = (\mathbf{I} - \frac{\Delta}{2}.\mathbf{A})^{-1} . \Delta.\mathbf{B} \quad (10)$$

- The discrete SSM converts the problem to a *sequence-to-sequence* mapping, useful for NLP.

Structured state-spaces

- The previous recurrent equation is not practical on modern hardware due to sequentiality.
- Instead, we write it as a **convolution**.
- If we enroll the discrete SSM :

$$x_k = \overline{\mathbf{A}^k \cdot \mathbf{B}}.u_0 + \overline{\mathbf{A}^{k-1} \cdot \mathbf{B}}.u_1 + \dots + \overline{\mathbf{B}}.u_k \quad (11)$$

$$y_k = \overline{\mathbf{C} \cdot \mathbf{A}^k \cdot \mathbf{B}}.u_0 + \overline{\mathbf{C} \cdot \mathbf{A}^{k-1} \cdot \mathbf{B}}.u_1 + \dots + \overline{\mathbf{C} \cdot \mathbf{B}}.u_k \quad (12)$$

$$(13)$$

- As a convolution :

$$y = \overline{\mathbf{K}} * u$$

with the kernel

$$\overline{\mathbf{K}} = (\overline{\mathbf{C} \cdot \mathbf{B}}, \overline{\mathbf{C} \cdot \mathbf{A} \cdot \mathbf{B}}, \dots, \overline{\mathbf{C} \cdot \mathbf{A}^{L-1} \cdot \mathbf{B}}) = (\overline{\mathbf{C} \cdot \mathbf{A}^i \cdot \mathbf{B}})_i$$

Structured state-spaces

- This was just the pre-requisite to the paper !
- We reduced the problem to a single convolution.
- It can be computed very efficiently using **Fast Fourier Transform** (FFT).
- *The main contribution of this paper is computing the SSM convolution kernel $\overline{\mathbf{K}}$ fast.*

Kernel computation

- Computing the kernel $\overline{\mathbf{K}}$ requires raising the matrix $\overline{\mathbf{A}}$ to several powers.
- That is fine if we diagonalize $\overline{\mathbf{A}}$, because conjugation is an equivalence relation in state-space models :

$$(\mathbf{A}, \mathbf{B}, \mathbf{C}) \sim (\mathbf{V}^{-1}.\mathbf{A}.\mathbf{V}, \mathbf{V}^{-1}.\mathbf{B}, \mathbf{C}.\mathbf{V}) \quad (14)$$

- The problem is that there exists a matrix \mathbf{V} diagonalizing the HiPPO matrix $\overline{\mathbf{A}}$, but it has entries exponentially large in the state space N ! Not stable numerically.

Diagonal plus low-rank

Instead of directly diagonalizing $\overline{\mathbf{A}}$, we decompose it into a **normal plus low-rank** matrix (linear algebra property) :

$$\mathbf{A} = \mathbf{V}.\mathbf{D}.\mathbf{V}^* - \mathbf{P}.\mathbf{Q}^T \quad (15)$$

which means conjugating to a *diagonal plus low-rank* (DPLR) :

$$\mathbf{A} = \mathbf{V} . (\mathbf{D} . - (\mathbf{V}^* \mathbf{P}) . (\mathbf{V}^* . \mathbf{Q})^*) . \mathbf{V}^* \quad (16)$$

The above is **Theorem 1** of the paper.

S4 Recurrence

Remember that we care about $\bar{\mathbf{A}}$, not \mathbf{A} :

$$\bar{\mathbf{A}} = (\mathbf{I} - \frac{\Delta}{2} \cdot \mathbf{A})^{-1} \cdot (\mathbf{I} + \frac{\Delta}{2} \cdot \mathbf{A}) \quad (17)$$

- Since \mathbf{A} is DPLR, so is the second term (obvious).
- The first term is also DPLR because of **Woodbury identity**.
- Thus, we obtain $\bar{\mathbf{A}}$ from \mathbf{A} in $O(N)$.
- In other words, one recurrent step of the discrete SSM has complexity $O(N)$ (**Theorem 2**).

S4 Convolution

- Instead of computing \mathbf{K} directly, we compute its spectrum by evaluating its truncated generating function $\sum_{j=0}^{L-1} \overline{\mathbf{K}}_j \zeta^j$ at the roots of unity ζ^j .
- $\overline{\mathbf{K}}$ is then retrieved with an inverse FFT.
- We can replace the matrix powers in the kernel with just an inverse :

$$\overline{\mathbf{K}}(z) = \sum_{i=1}^{L-1} \overline{\mathbf{C} \cdot \mathbf{A}^i \cdot \mathbf{B}} \cdot z^i = \dots = \tilde{\mathbf{C}} \cdot (\mathbf{I} - \overline{\mathbf{A}} \cdot z)^{-1} \cdot \overline{\mathbf{B}} \quad (18)$$

- However, this inverse needs to be calculated L times (once for each root of the unity).

S4 Convolution

- Evaluating the kernel $\overline{\mathbf{A}}$ to the roots of the unity, because of the special structure of the matrix \mathbf{A} that we explored before, ends up becoming a **Cauchy kernel**.
- A Cauchy kernel can be evaluated efficiently in $O(N + L) \cdot \log^2(N + L)$ operations (while it would be $O(N \cdot L)$ naively).
- The S4 convolutional kernel involves 4 Cauchy kernels, requiring $O(N + L)$ operations and $O(N + L)$ space (**Theorem 3**).

S4 Layer

- One S4 layer has the following learnable parameters : **D**, **P**, **Q**, **B** and **C**.
- An S4 layer maps a \mathbb{R}^L signal to another \mathbb{R}^L signal.
- We make it process H features instead of 1 :
 - Define H independant copies of the S4 layer.
 - Mix the features with a position-wise linear layer.
 - Add non-linear activations between these layers.
- After this process, S4 defines a sequence-to-sequence map of shape (batch size, sequence length, hidden dimension), just like Transformer.

Algorithm

Overview of S4 algorithm to compute the convolutional kernel :

Algorithm 1 S4 CONVOLUTION KERNEL (SKETCH)

Input: S4 parameters $\Lambda, P, Q, B, C \in \mathbb{C}^N$ and step size Δ

Output: SSM convolution kernel $\bar{\mathbf{K}} = \mathcal{K}_L(\bar{\mathbf{A}}, \bar{\mathbf{B}}, \bar{\mathbf{C}})$ for $\mathbf{A} = \mathbf{\Lambda} - \mathbf{P}\mathbf{Q}^*$ (equation (5))

- 1: $\tilde{\mathbf{C}} \leftarrow (\mathbf{I} - \overline{\mathbf{A}}^L)^* \overline{\mathbf{C}}$ ▷ Truncate SSM generating function (SSMGF) to length L
- 2: $\begin{bmatrix} k_{00}(\omega) & k_{01}(\omega) \\ k_{10}(\omega) & k_{11}(\omega) \end{bmatrix} \leftarrow [\tilde{\mathbf{C}} \mathbf{Q}]^* \left(\frac{2}{\Delta} \frac{1-\omega}{1+\omega} - \mathbf{\Lambda} \right)^{-1} [\mathbf{B} \mathbf{P}]$ ▷ Black-box Cauchy kernel
- 3: $\hat{\mathbf{K}}(\omega) \leftarrow \frac{2}{1+\omega} [k_{00}(\omega) - k_{01}(\omega)(1 + k_{11}(\omega))^{-1} k_{10}(\omega)]$ ▷ Woodbury Identity
- 4: $\hat{\mathbf{K}} = \{\hat{\mathbf{K}}(\omega) : \omega = \exp(2\pi i \frac{k}{L})\}$ ▷ Evaluate SSMGF at all roots of unity $\omega \in \Omega_L$
- 5: $\overline{\mathbf{K}} \leftarrow \text{iFFT}(\hat{\mathbf{K}})$ ▷ Inverse Fourier Transform

Complexity

L is sequence length, B is the batch size, H is the feature size.

	Convolution ³	Recurrence	Attention	S4
Parameters	LH	H^2	H^2	H^2
Training	$\tilde{L}H(\mathbf{B} + \mathbf{H})$	BLH^2	$B(L^2H + LH^2)$	$\mathbf{B}H(\tilde{H} + \tilde{L}) + B\tilde{L}H$
Space	BLH	BLH	$B(L^2 + HL)$	BLH
Parallel	Yes	No	Yes	Yes
Inference	LH^2	H^2	$L^2H + H^2L$	H^2

S4 is optimal or near-optimal on all aspects.

Long-Range Arena

LRA is a benchmark of several tasks with very long-range dependencies :

MODEL	LISTOPS	TEXT	RETRIEVAL	IMAGE	PATHFINDER	PATH-X	AVG
Transformer	36.37	64.27	57.46	42.44	71.40	✗	53.66
Reformer	<u>37.27</u>	56.10	53.40	38.07	68.50	✗	50.56
BigBird	36.05	64.02	59.29	40.83	74.87	✗	54.17
Linear Trans.	16.13	<u>65.90</u>	53.09	42.34	75.30	✗	50.46
Performer	18.01	65.40	53.82	42.77	77.05	✗	51.18
FNNet	35.33	65.11	59.61	38.67	<u>77.80</u>	✗	54.42
Nyströmformer	37.15	65.52	<u>79.56</u>	41.58	<u>70.94</u>	✗	57.46
Luna-256	37.25	64.57	79.29	<u>47.38</u>	77.72	✗	<u>59.37</u>
S4	58.35	76.02	87.09	87.26	86.05	88.10	80.48

S4 is SOTA by a very large margin, and is the first model ever to beat random chance on Path-X (Path-X is a Pathfinder task, aka identifying the correct target contour of a shape, but with very high resolution (16k) images).

Speech classification

Speech classification accuracy with standard MFCC features, unprocessed signals (Raw), and frequency change at test time (x0.5).

	MFCC	RAW	0.5×
Transformer	90.75	X	X
Performer	80.85	30.77	30.68
ODE-RNN	65.9	X	X
NRDE	89.8	16.49	15.12
ExpRNN	82.13	11.6	10.8
LipschitzRNN	88.38	X	X
CKConv	95.3	71.66	<u>65.96</u>
WaveGAN-D	X	<u>96.25</u>	X
LSSL	93.58	X	X
S4	<u>93.96</u>	98.32	96.30

Image classification

Pixel-level sequential image classification, sliding non-overlapping windows of 1024 pixels.

	sMNIST	pMNIST	sCIFAR
Transformer	98.9	97.9	62.2
LSTM	98.9	95.11	63.01
r-LSTM	98.4	95.2	72.2
UR-LSTM	99.28	96.96	71.00
UR-GRU	99.27	96.51	74.4
HiPPO-RNN	98.9	98.3	61.1
LMU-FFT	-	98.49	-
LipschitzRNN	99.4	96.3	64.2
TCN	99.0	97.2	-
TrellisNet	99.20	98.13	73.42
CKConv	99.32	98.54	63.74
LSSL	<u>99.53</u>	98.76	<u>84.65</u>
S4	99.63	<u>98.70</u>	91.13

S4 is better than Transformer, LSTMs, and CNNs.

Image density estimation

Density estimation on CIFAR-10 :

Model	bpd	2D bias	Images / sec
Transformer	3.47	None	0.32 ($1\times$)
Linear Transf.	3.40	None	17.85 ($56\times$)
PixelCNN	3.14	2D conv.	-
Row PixelRNN	3.00	2D BiLSTM	-
PixelCNN++	2.92	2D conv.	<u>19.19</u> ($59.97\times$)
Image Transf.	2.90	2D local attn.	0.54 ($1.7\times$)
PixelSNAIL	<u>2.85</u>	2D conv. + attn.	0.13 ($0.4\times$)
Sparse Transf.	2.80	2D sparse attn.	-
S4 (base)	2.92	None	20.84 ($65.1\times$)
S4 (large)	<u>2.85</u>	None	3.36 ($10.5\times$)

S4 is competitive with SOTA in terms of bits per dimension.
S4-base is the fastest overall.

Language modeling

Word-level language modeling (perplexity) on WikiText-103 :

Model	Params	Test ppl.	Tokens / sec
Transformer	247M	20.51	0.8K (1×)
GLU CNN	229M	37.2	-
AWD-QRNN	151M	33.0	-
LSTM + Hebb.	-	29.2	-
TrellisNet	180M	29.19	-
Dynamic Conv.	255M	25.0	-
TaLK Conv.	240M	23.3	-
S4	249M	21.28	48K (60×)

A Transformer with attention blocks replaced by S4 blocks provides the best perplexity among attention-free models, very close to Transformer with self-attention.

Besides S4 processes 60x more tokens per second.

Conclusion

- A fresh departure from the classical Transformer layers and self-attention.
- Very high performing on long sequence modeling.
- Also computationnally efficient, as it's not limited by the $O(L^2)$ complexity of self-attention.
- A method derived from an old tool (state-space models), rendered useful through applying multiple linear algebra tricks.
- I am currently working with an intern at A*STAR on how to extend S4 to an encoder-decoder model, with the goal of applying it to long-input summarization (e.g, Arxiv, PubMed), perhaps document-level translation too.